

Machine learning in science and society

From automated science to beneficial artificial intelligence

Christos Dimitrakakis

May 30, 2019

Introduction

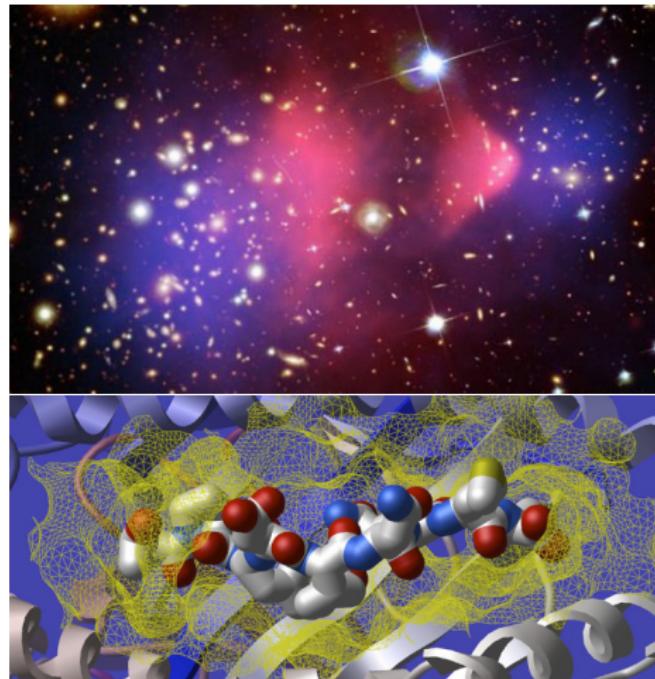
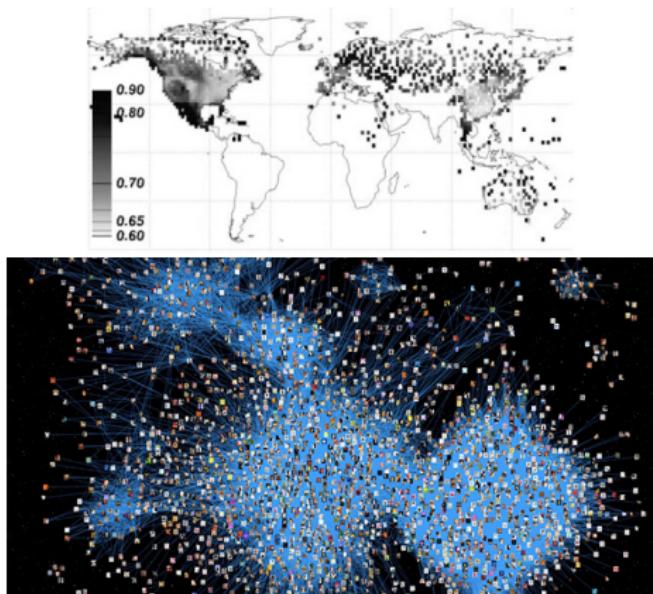
1 Introduction to machine learning

- Data analysis, learning and planning
- Experiment design
- Bayesian inference.
- Course overview

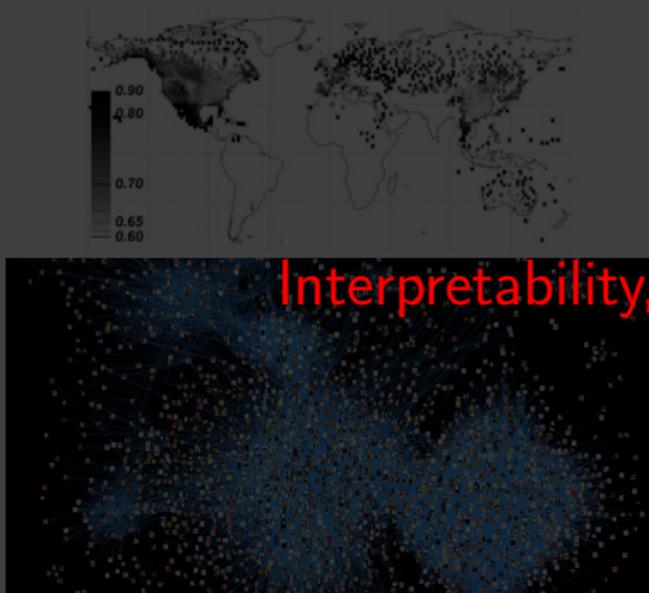
2 Nearest neighbours

3 Reproducibility

Scientific applications



Scientific applications



Interpretability,



Reproducibility

Pervasive “intelligent” systems



Home assistants



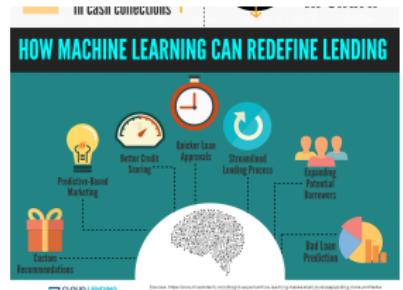
Autonomous vehicles



Web advertising



Ridesharing



Lending



Public policy

Pervasive “intelligent” systems



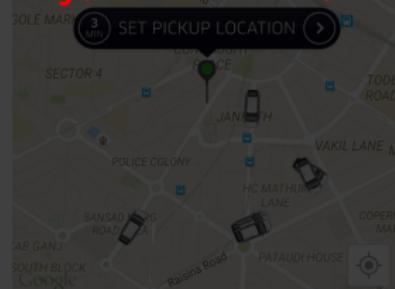
Home assistants



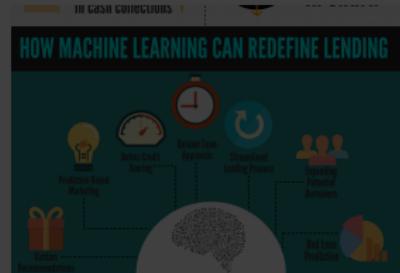
Web advertising



Autonomous vehicles



Ridesharing



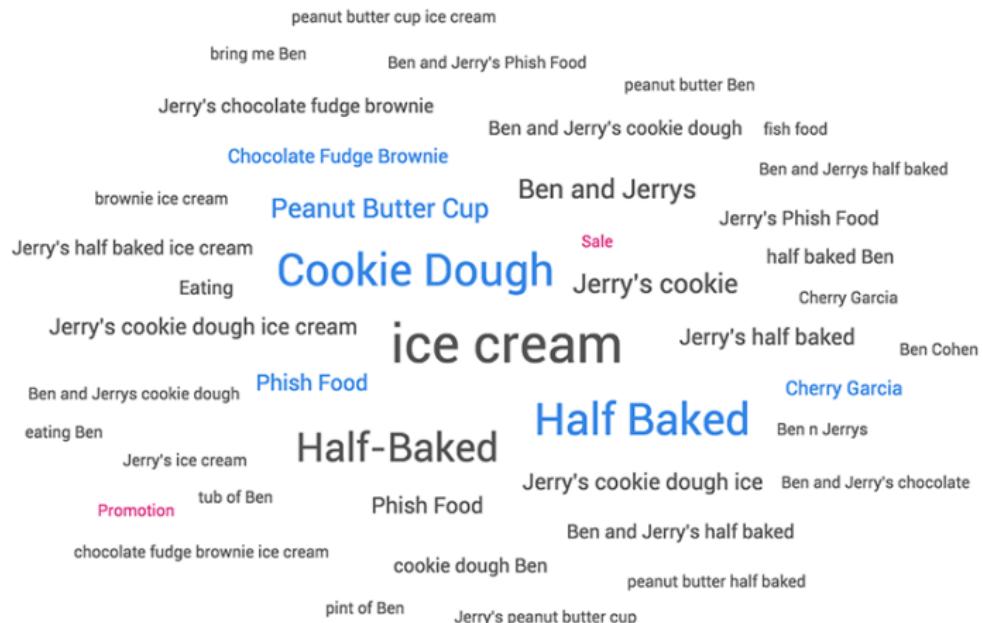
Lending



Public policy

What can machine learning do?

Can machines learn from data?



■ Topics ■ Tags ■ Categories

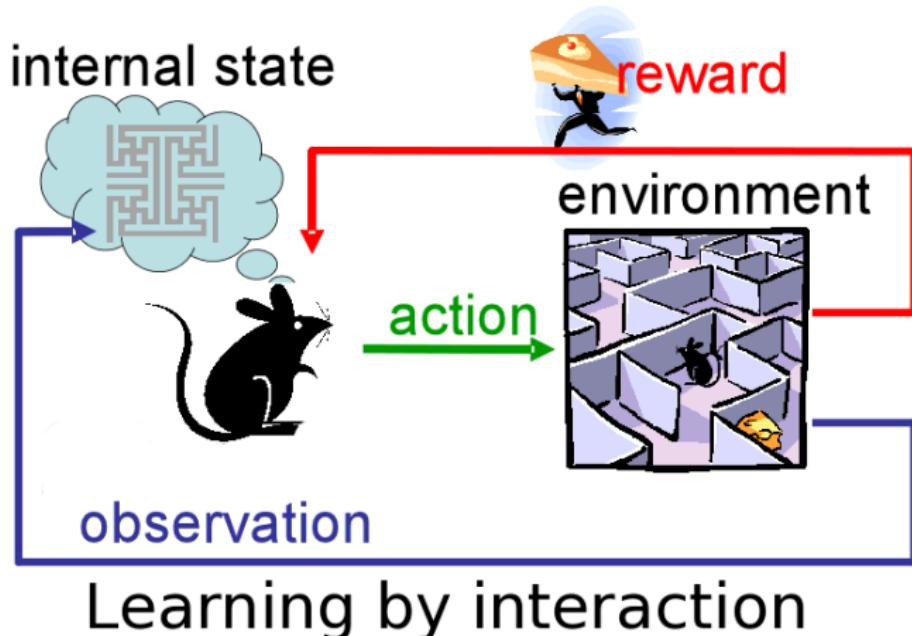
An unsupervised learning problem: topic modelling

Can machines learn from data?



A supervised learning problem: object recognition

Can machines learn from their mistakes?



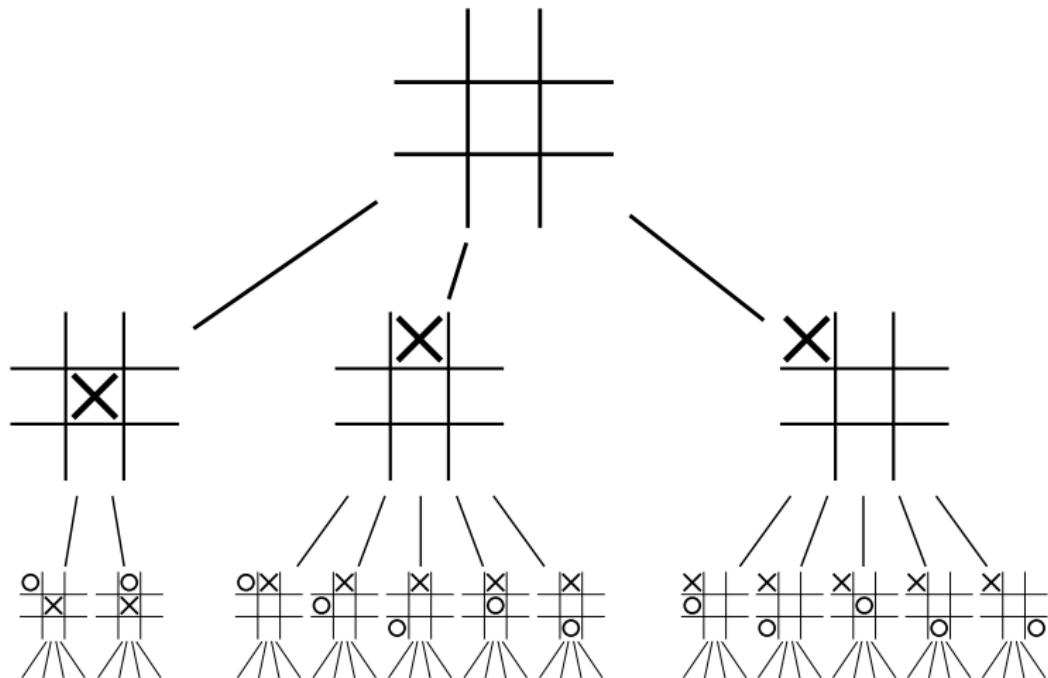
Reinforcement learning

Take actions a_1, \dots, a_t , so as to maximise utility $U = \sum_{t=1}^T r_t$

Can machines make complex plans?



Machines can make complex plans!



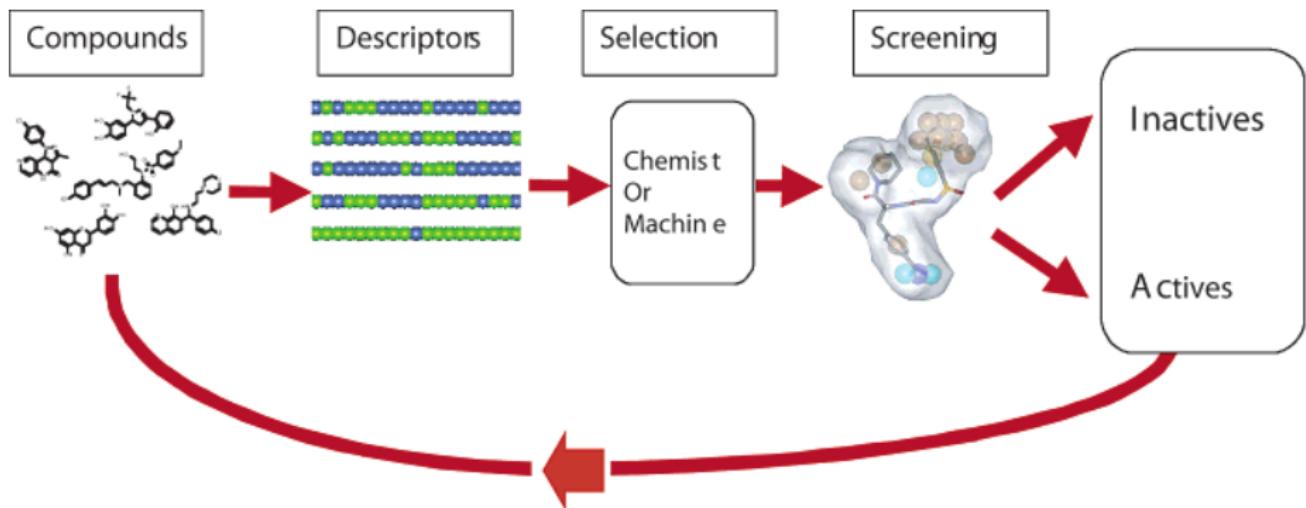
The scientific process as machine learning



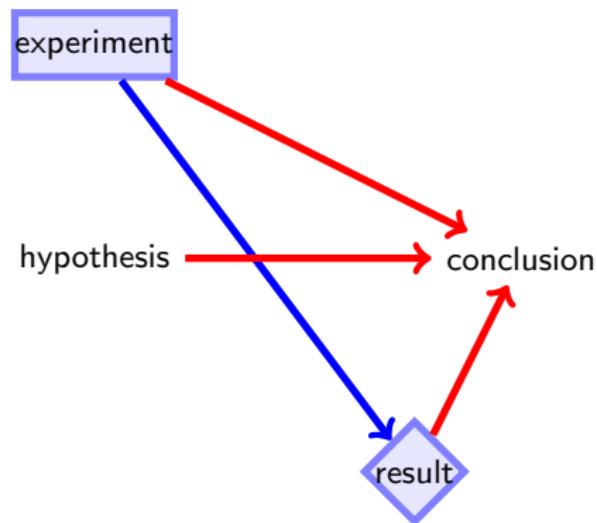
Adam, the robot scientist



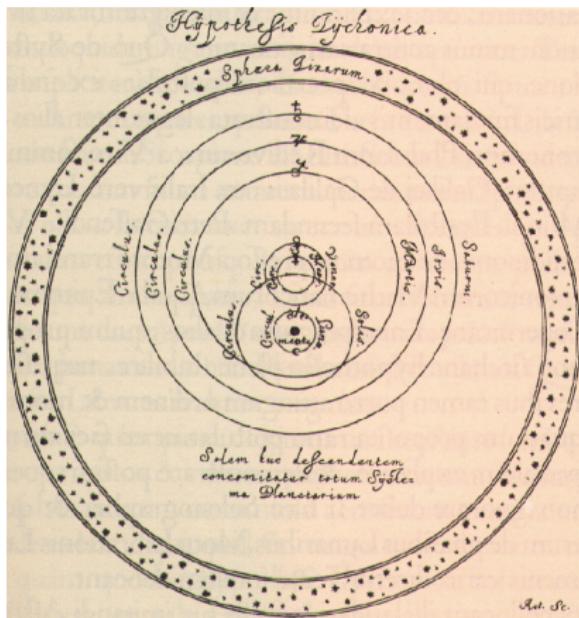
Drug discovery



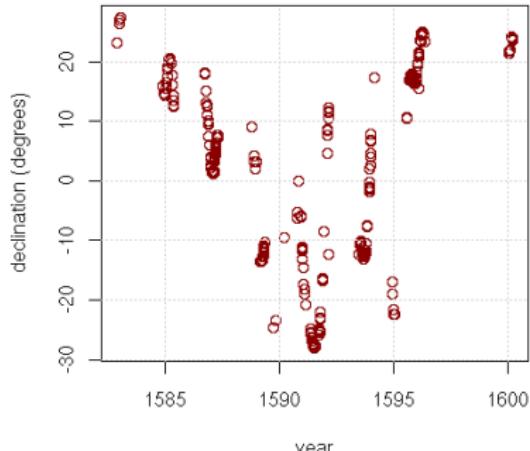
Drawing conclusions from results



Tycho Brahe's minute eye measurements



Tycho Brahe's Mars Observations



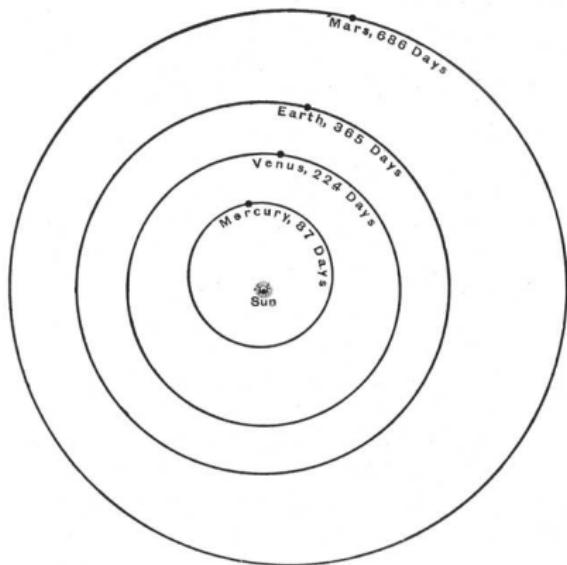
source: Tychonis Brahe Dani Opera Omnia

Figure: Tycho's measurements of the orbit of Mars and the conclusion about the actual orbits, under the assumption of an earth-centric universe with circular orbits.

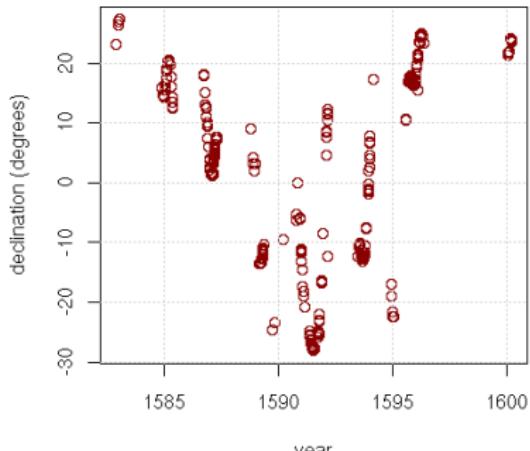
- Hypothesis: Earth-centric, Circular orbits
- Conclusion: Specific circular orbits

C. Dimitrakakis

Johannes Kepler's alternative hypothesis



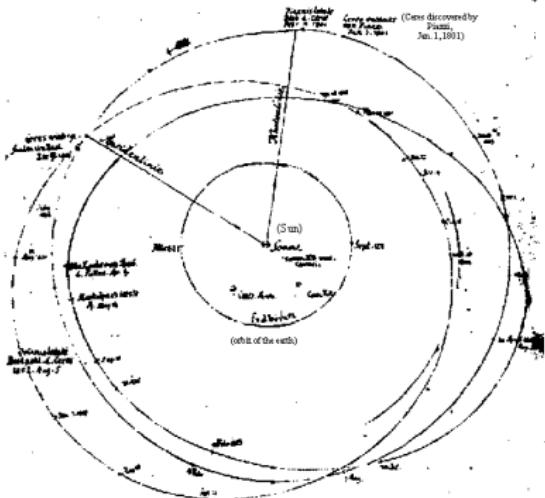
Tycho Brahe's Mars Observations



source: Tychonis Brahe Dani Opera Omnia

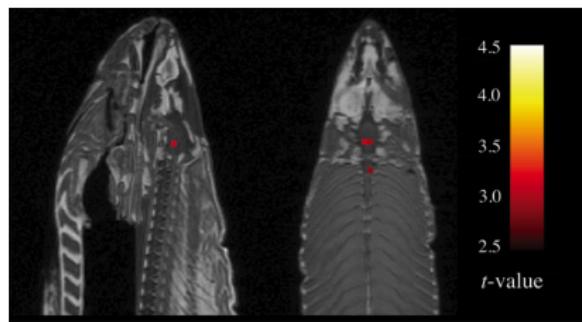
- Hypothesis: Circular or elliptic orbits
- Conclusion: Specific elliptic orbits

200 years later, Gauss formalised this statistically

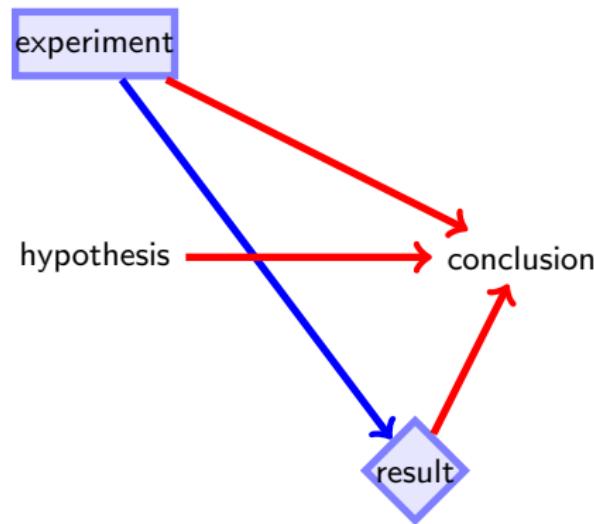


Sketch of the orbits of Ceres and Pallas (nachlaß Gauß, Handb. 4). Courtesy of Universitätsbibliothek Göttingen.

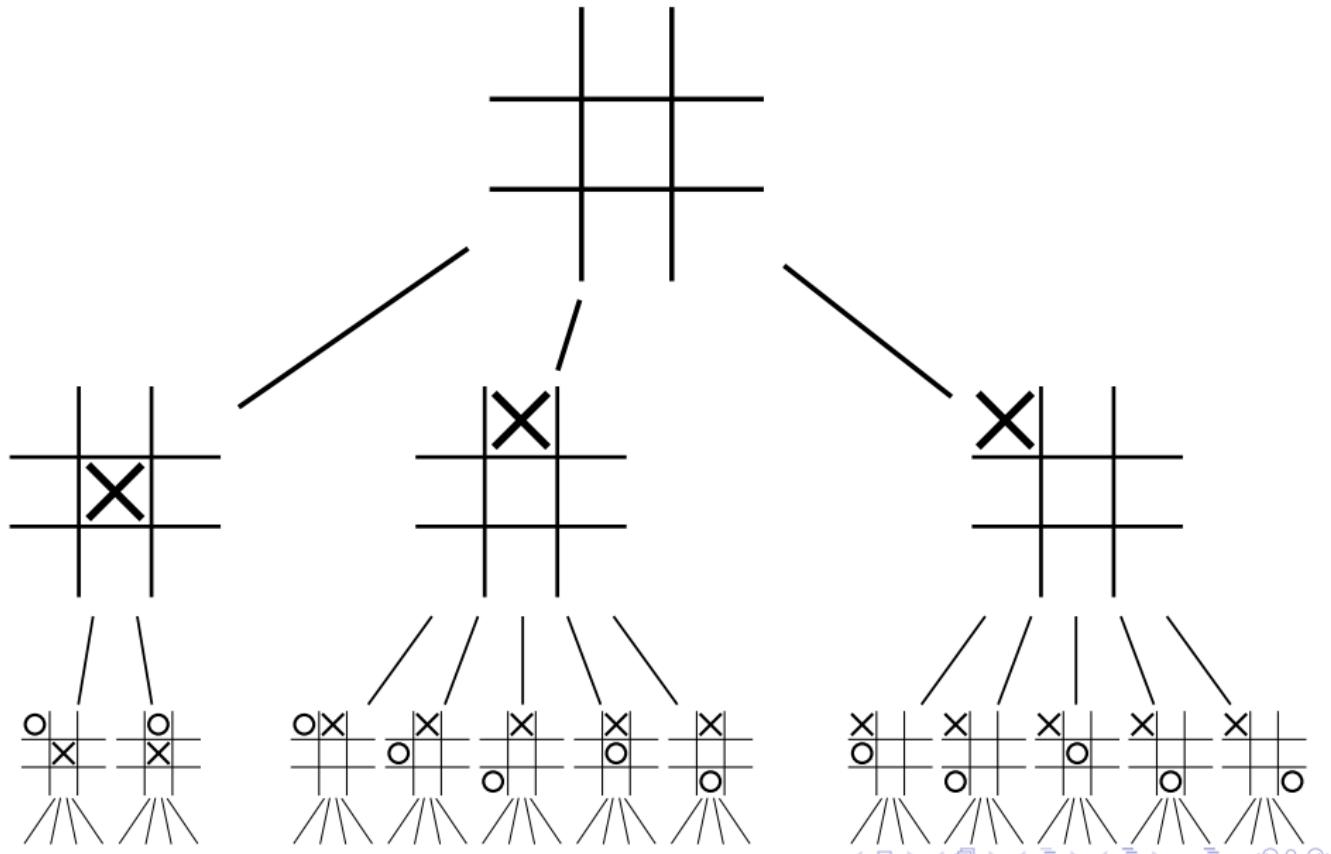
A warning: The dead salmon mirage



Planning future experiments



Planning experiments is like Tic-Tac-Toe



Eve, another robot scientist



a malaria drug

Machine learning in practice

Avoiding pitfalls

- Choosing hypotheses.
- Correctly interpreting conclusions.
- Using a good testing methodology.

Machine learning in society

- Privacy
- Fairness
- Safety

Machine learning in practice

Avoiding pitfalls

- Choosing hypotheses.
- Correctly interpreting conclusions.
- Using a good testing methodology.

Machine learning in society

- Privacy — Credit risk.
- Fairness
- Safety

Machine learning in practice

Avoiding pitfalls

- Choosing hypotheses.
- Correctly interpreting conclusions.
- Using a good testing methodology.

Machine learning in society

- Privacy — Credit risk.
- Fairness — Job market.
- Safety

Machine learning in practice

Avoiding pitfalls

- Choosing hypotheses.
- Correctly interpreting conclusions.
- Using a good testing methodology.

Machine learning in society

- Privacy — Credit risk.
- Fairness — Job market.
- Safety — Medicine.

Technical topics

Machine learning problems

- Unsupervised learning.
- Supervised learning.
- Reinforcement learning.

Algorithms and models

- Bayesian inference and graphical models.
- Stochastic optimisation and neural networks.
- Backwards induction and Markov decision processes.

Course structure

Module structure

- **Activity**-based, hands-on.
- Mini-lectures with short exercises in each class.
- Technical tutorials and labs in alternate week.

Modules

Three mini-projects.

- Simple decision problems: Credit risk.
- Structured problems: Fake news.
- Sequential problems: Medical diagnostics and treatment.

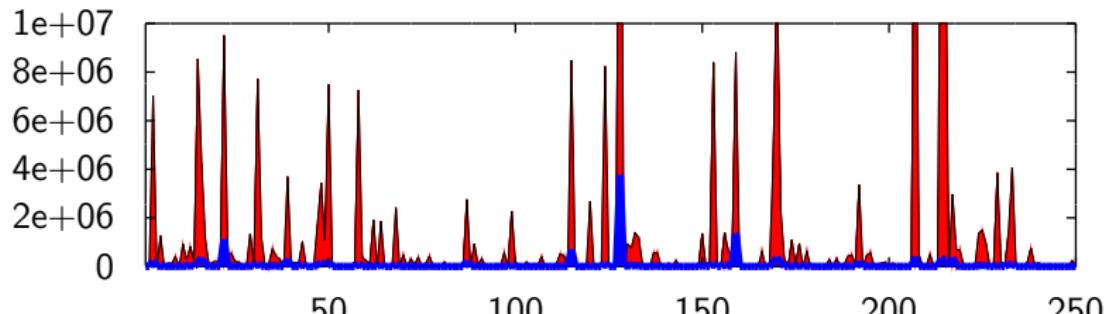
1 Introduction to machine learning

2 Nearest neighbours

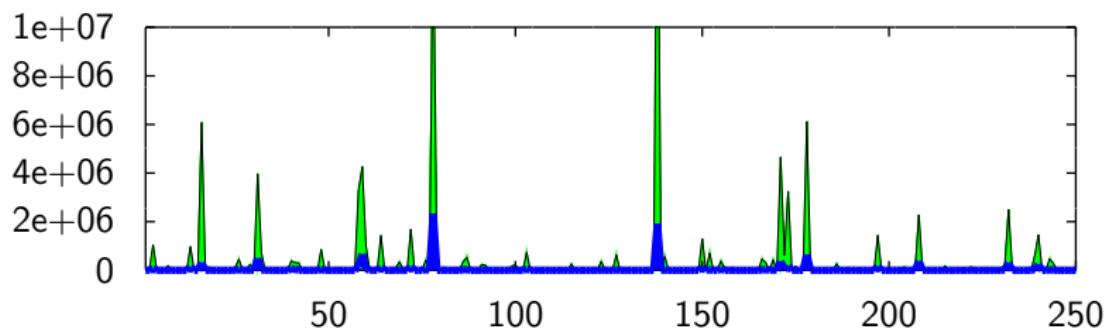
3 Reproducibility

Discriminating between diseases

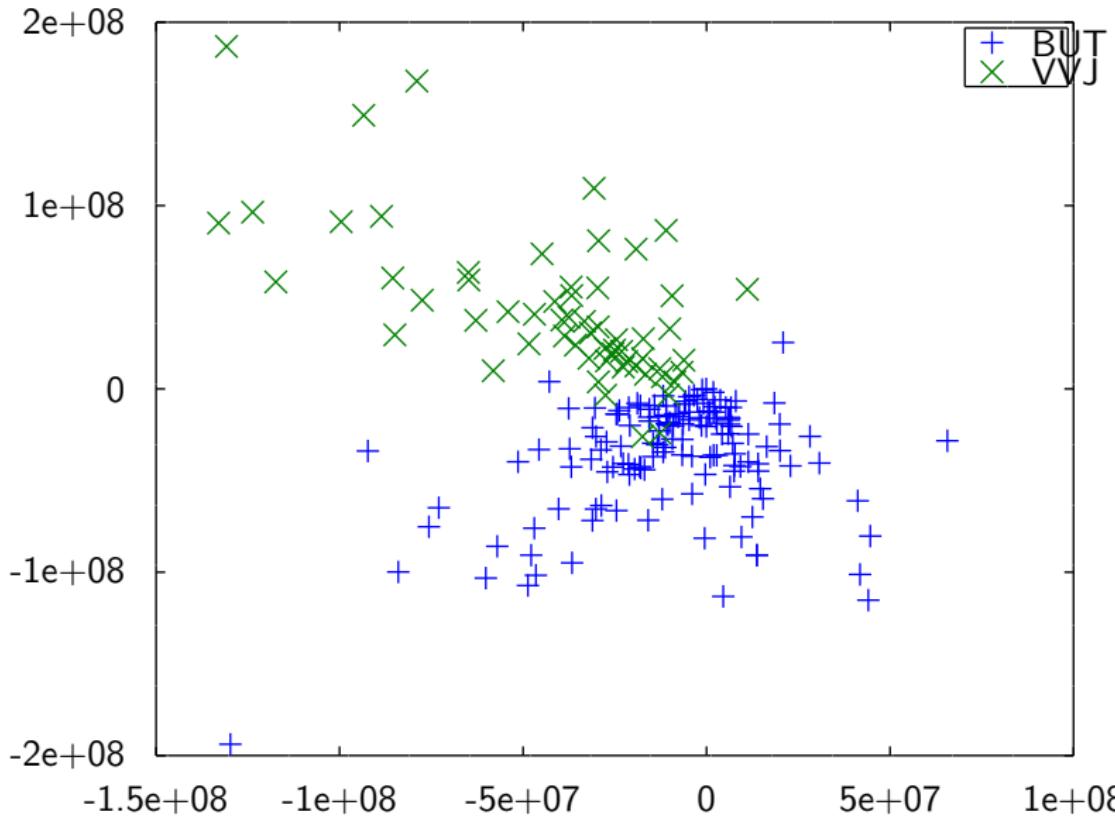
Spectral statistics VVX strain



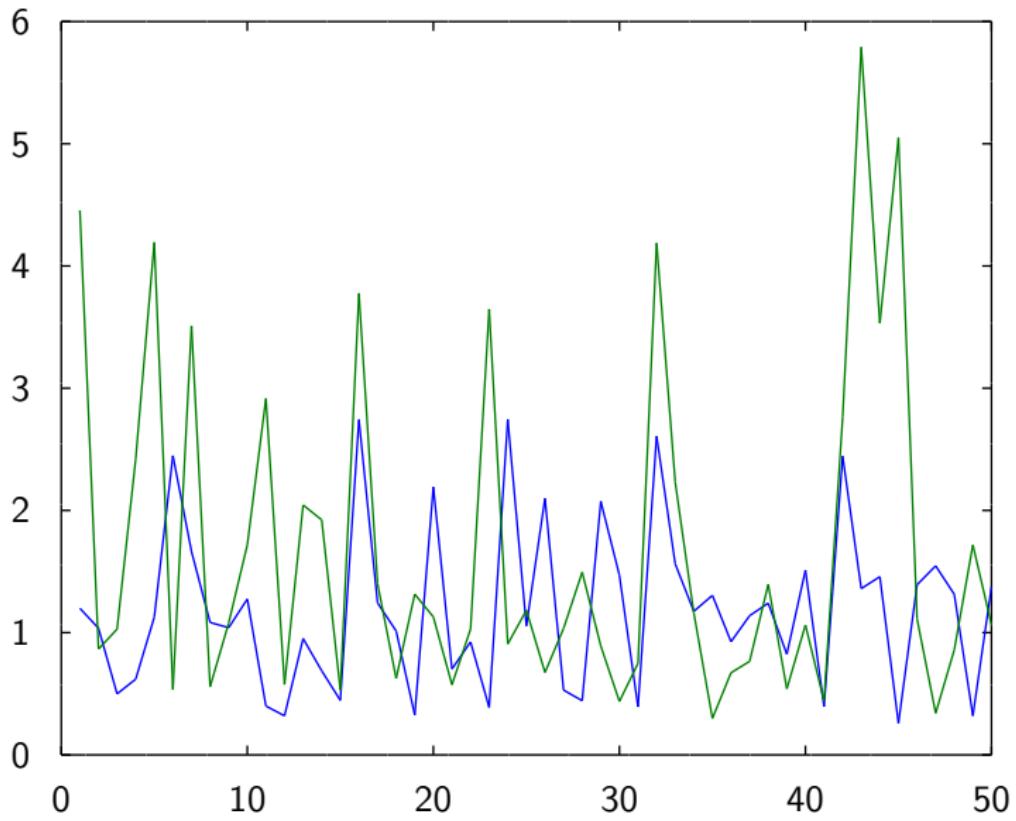
Spectral statistics for BUT



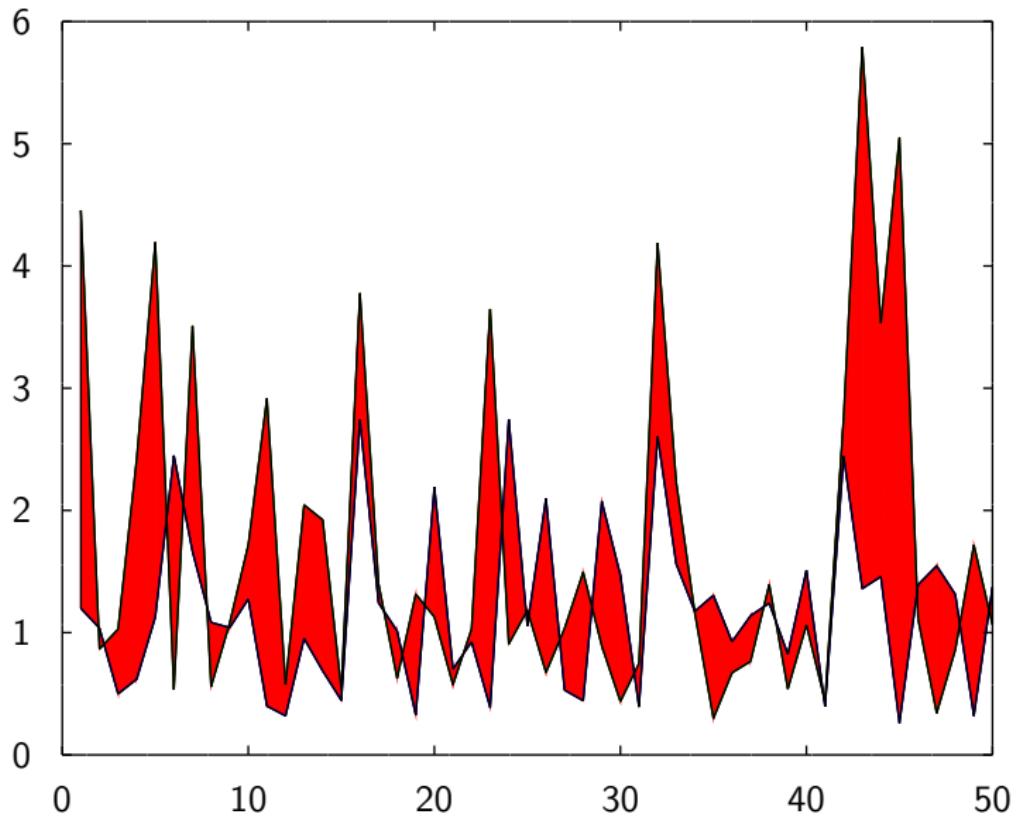
Nearest neighbour: the hidden secret of machine learning



Comparing spectral data



Comparing spectral data



The nearest neighbour algorithm

Algorithm 1 k -NN Classify

- 1: **Input** Data $D = \{(x_1, y_1), \dots, (x_T, y_T)\}$, $k \geq 1$, $d : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}_+$, new point $x \in \mathcal{X}$
 - 2: $D = \text{Sort}(D, d)$ % Sort D so that $d(x, x_i) \leq d(x, x_{i+1})$.
 - 3: $p_y = \sum_{i=1}^k \mathbb{I}\{y_i = y\} / k$ for $y \in \mathcal{Y}$.
 - 4: **Return** $p \triangleq (p_1, \dots, p_k)$
-

Algorithm parameters

- Neighbourhood $k \geq 1$.
- Distance $d : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}_+$.

What does the algorithm output when $k = T$?

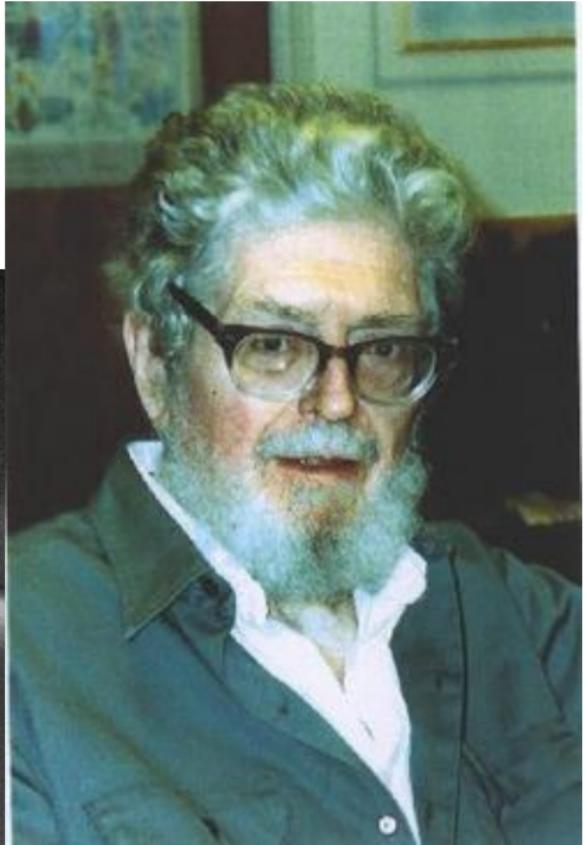
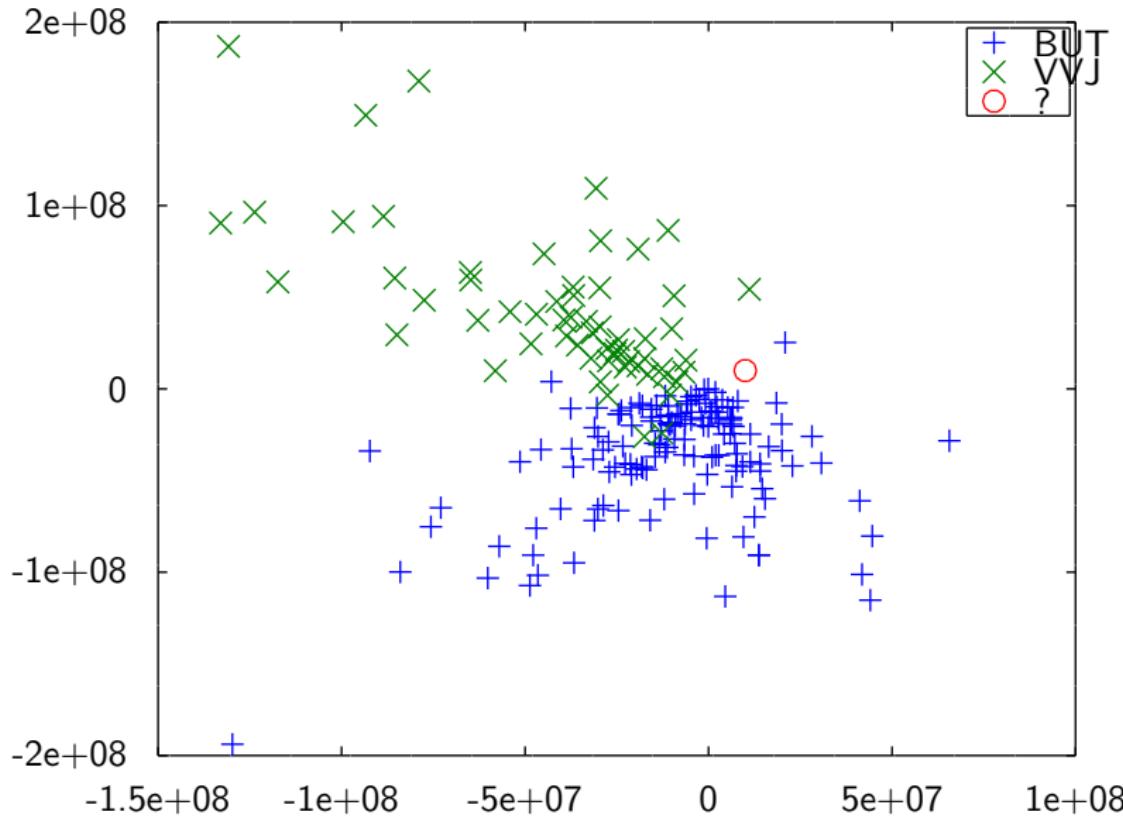
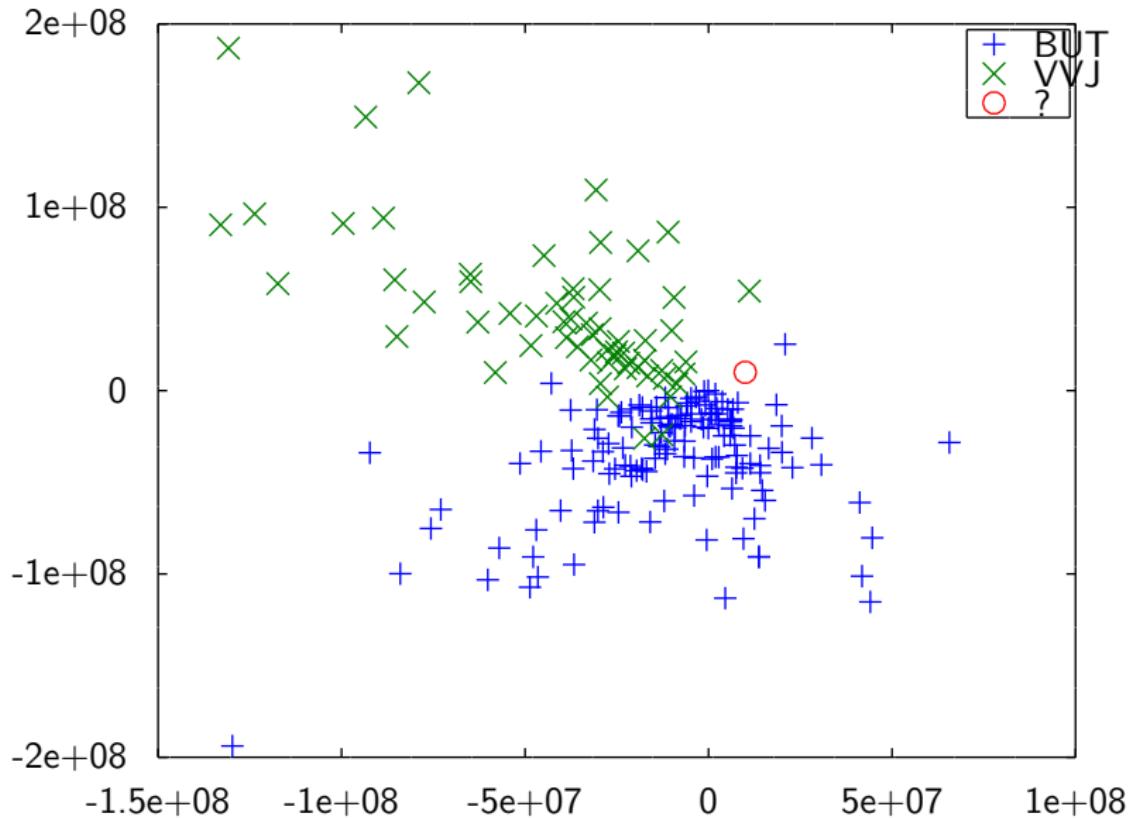


Figure: The nearest neighbours algorithm was introduced by Fix and Hodges Jr¹, who also proved consistency properties.

Nearest neighbour: What type is the new bacterium?



Nearest neighbour: What type is the new bacterium?



What if it's a completely different strain?

Separating the model from the classification policy

- The k -NN algorithm returns a model giving class probabilities for new data points.
- Deciding a class given the model

$$\pi(a | x) = \mathbb{I}\{p_a \geq p_y \forall y\}, \quad p = k\text{-NN}(D, k, d, x)$$

Hands on with Python console

KNN example

Discussion: Shortcomings of k -nearest neighbour

- Choice of k
- Choice of metric d .
- Representation of uncertainty.
- Scaling with large amounts of data.
- Meaning of label probabilities.

Learning outcomes

Understanding

- How kNN works
- The effect of hyperparameters k, d for nearest neighbour.
- The use of kNN to classify new data.

Skills

- Use a standard kNN class in python
- Optimise kNN hyperparameters in an unbiased manner.
- Calculate probabilities of class labels using kNN.

Reflection

- When is kNN a good model?
- How can we deal with large amounts of data?
- How can we best represent uncertainty?

1 Introduction to machine learning

2 Nearest neighbours

3 Reproducibility

- The human as an algorithm
- Algorithmic sensitivity

Computational reproducibility: Can the study be repeated?

Can we, from the available information and data, exactly reproduce the reported methods and results?

- jupyter notebooks
- svn, git or mercurial version control systems

Scientific reproducibility: Is the conclusion correct?

Can we, from the available information and a **new** set of data, reproduce the conclusions of the original study?

When publishing results about a **new method**, computational reproducibility is essential for scientific reproducibility.

RealClear Politics

[Polls](#) [Election 2018](#) [Video](#) [Writers](#) [More](#)

Poll	Date	Sample	MoE	Clinton (D)	Trump (R)	Spread
Final Results	--	--	--	48.2	46.1	Clinton +2.1
RCP Average	11/1 - 11/7	--	--	46.8	43.6	Clinton +3.2
Bloomberg	11/4 - 11/6	799 LV	3.5	46	43	Clinton +3
IBD/TIPP Tracking	11/4 - 11/7	1107 LV	3.1	43	42	Clinton +1
Economist/YouGov	11/4 - 11/7	3669 LV	-	49	45	Clinton +4
LA Times/USC Tracking	11/1 - 11/7	2935 LV	4.5	44	47	Trump +3
ABC/Wash Post Tracking	11/3 - 11/6	2220 LV	2.5	49	46	Clinton +3
FOX News	11/3 - 11/6	1295 LV	2.5	48	44	Clinton +4
Monmouth	11/3 - 11/6	748 LV	3.6	50	44	Clinton +6
NBC News/Wall St. Jrnal	11/3 - 11/5	1282 LV	2.7	48	43	Clinton +5
CBS News	11/2 - 11/6	1426 LV	3.0	47	43	Clinton +4
Reuters/Ipsos	11/2 - 11/6	2196 LV	2.3	44	39	Clinton +5

All General Election: Trump vs. Clinton Polling Data



RCP POLL AVERAGE

General Election: Trump vs.
Clinton

46.8	Clinton (D) +3.2
43.6	Trump (R)



The principle of independent evaluation

Data used for estimation cannot be used for evaluation.

Data Collection



Figure: The decision process in classification.



Figure: The decision process in classification.



Algorithm, hyperparameters

Figure: The decision process in classification.

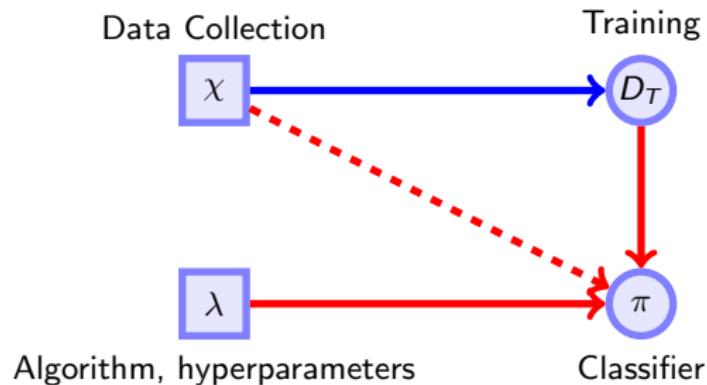


Figure: The decision process in classification.

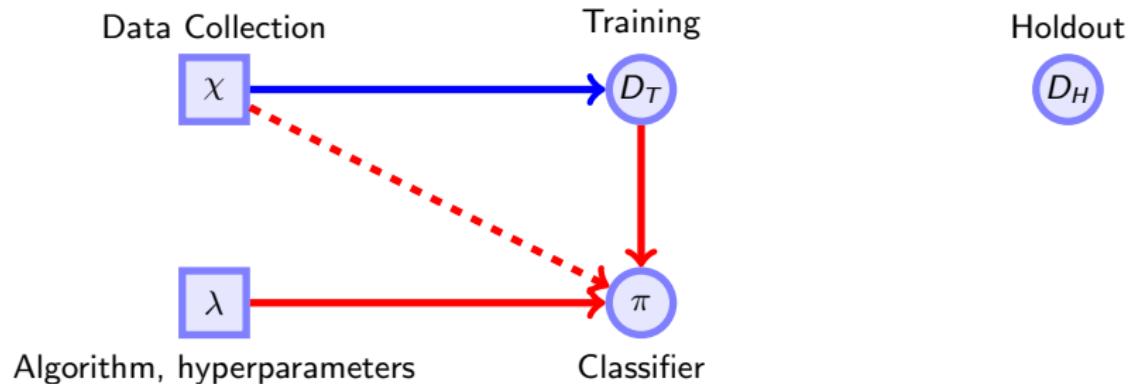


Figure: The decision process in classification.

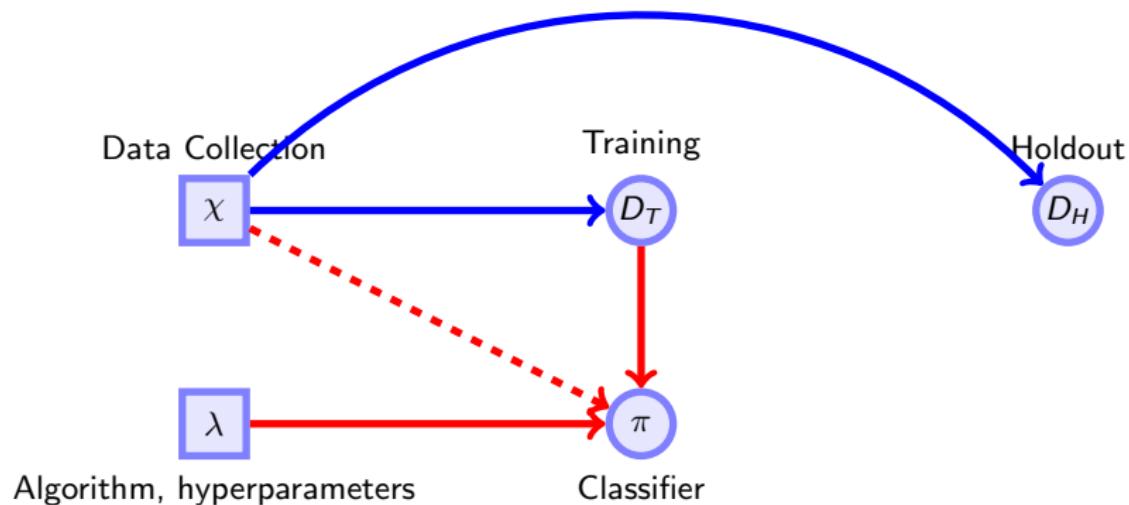


Figure: The decision process in classification.

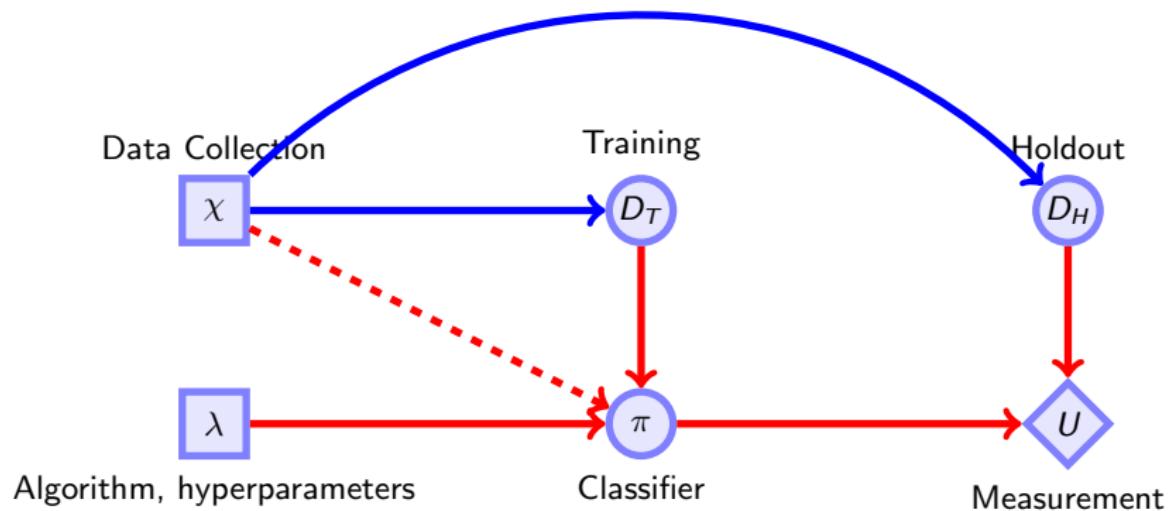


Figure: The decision process in classification.

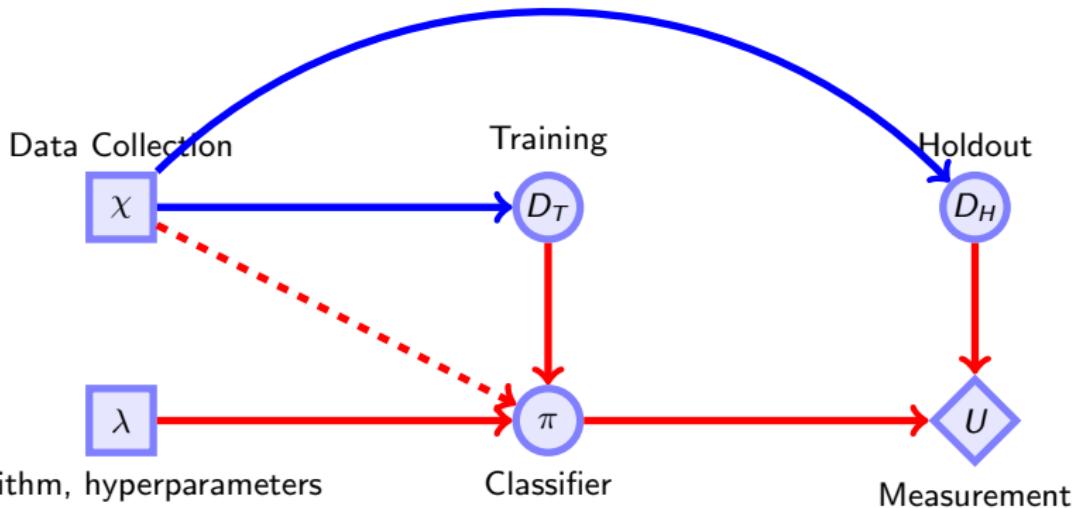


Figure: The decision process in classification.

Classification accuracy

$$\mathbb{E}_x[U(\pi)] = \sum_{x,y} \underbrace{\mathbb{P}_x(x,y)}_{\text{Data probability}} \overbrace{\pi(a=y | x)}^{\text{Decision probability}}$$

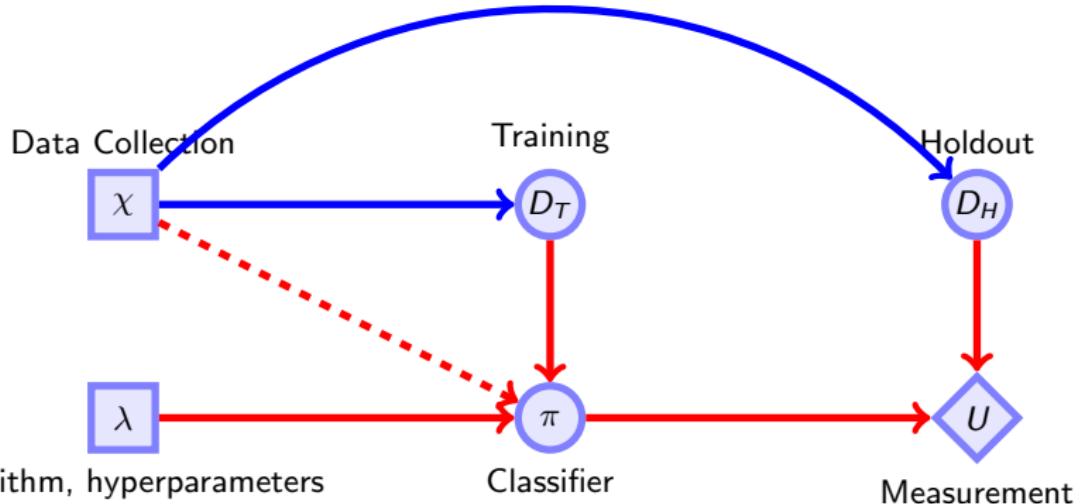


Figure: The decision process in classification.

Classification accuracy

$$\mathbb{E}_{D_H} U(\pi) = \sum_{(x,y) \in D_H} \pi(a = y | x) / |D_H|.$$

The human as an algorithm.

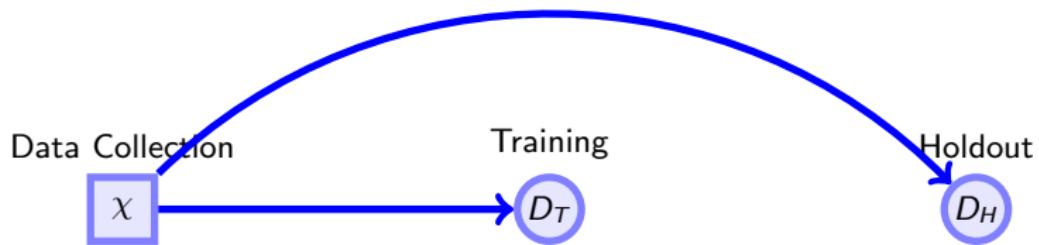


Figure: Selecting algorithms and hyperparameters through holdouts

The human as an algorithm.

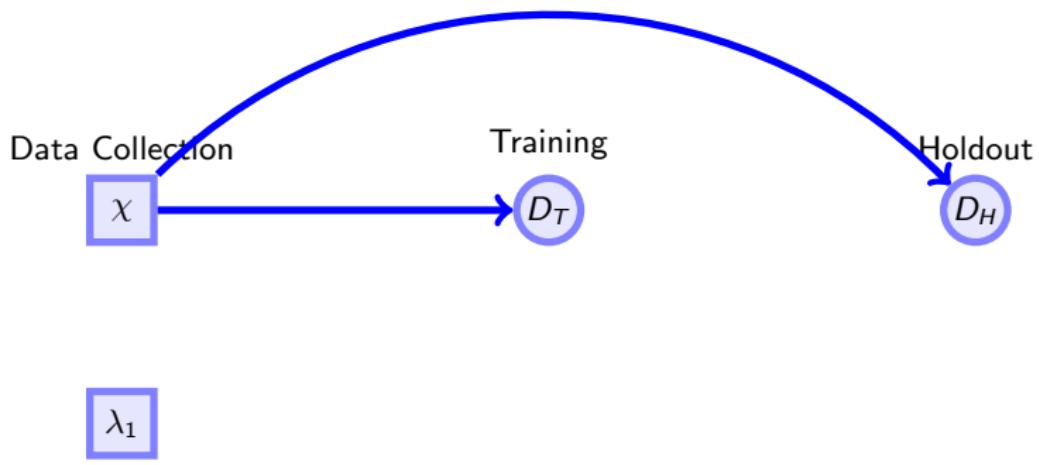


Figure: Selecting algorithms and hyperparameters through holdouts

The human as an algorithm.

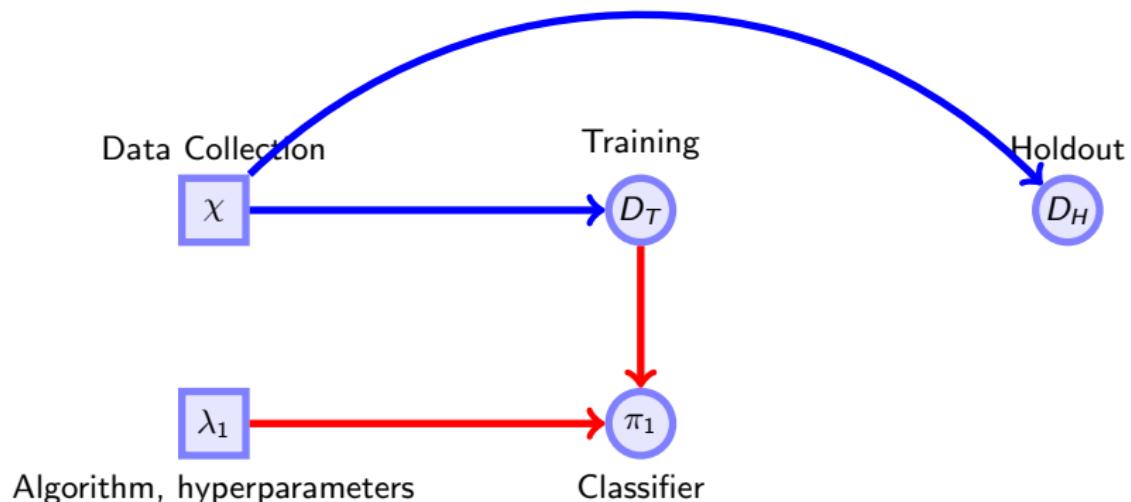


Figure: Selecting algorithms and hyperparameters through holdouts

The human as an algorithm.

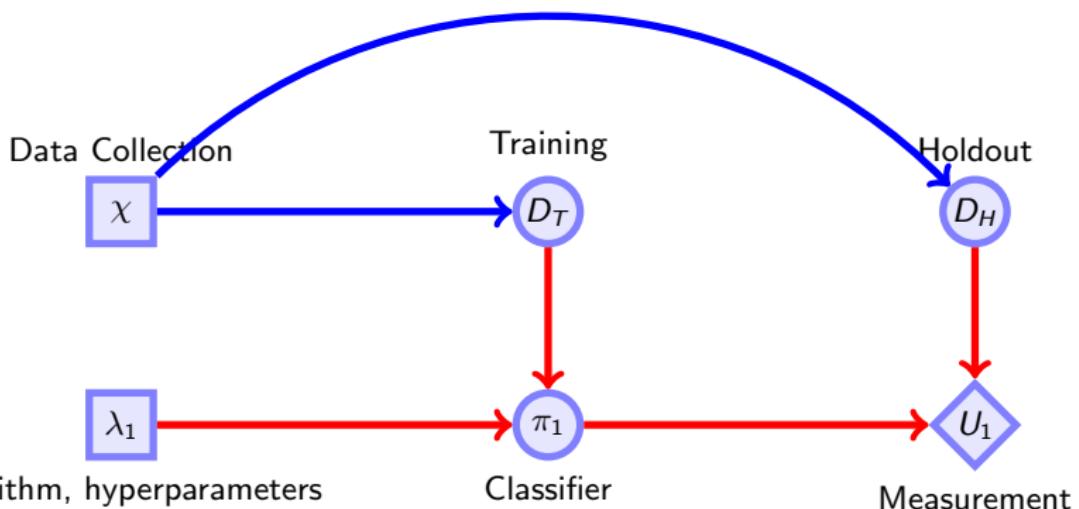


Figure: Selecting algorithms and hyperparameters through holdouts

The human as an algorithm.

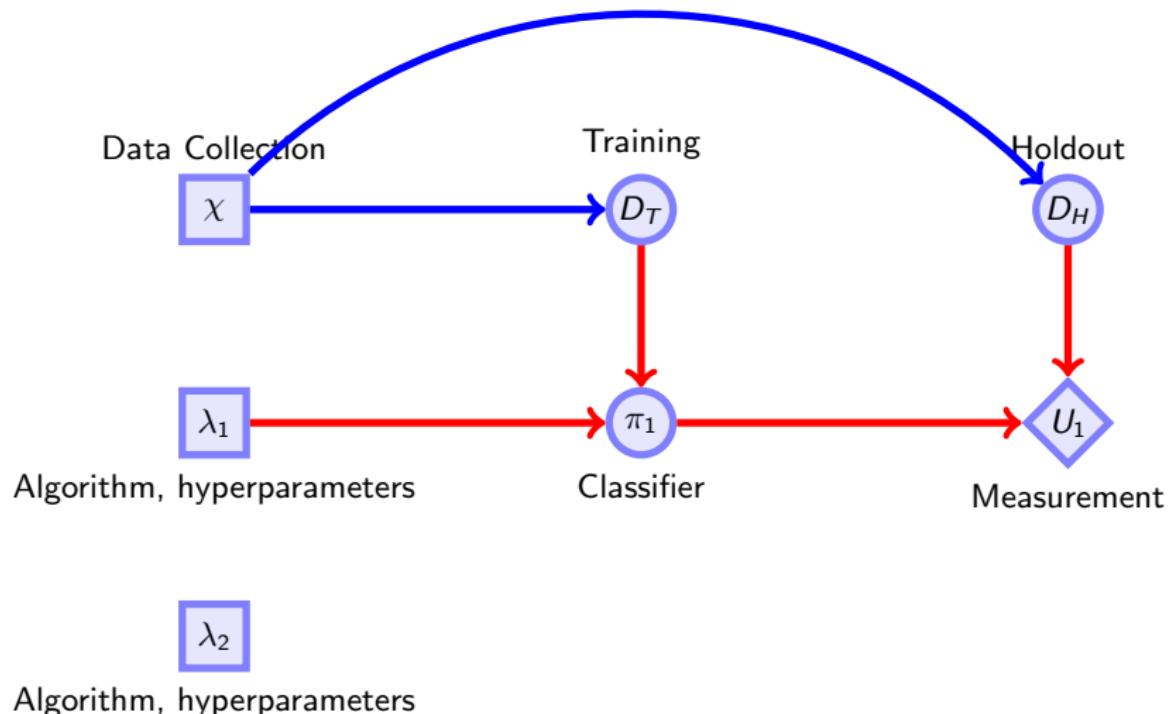


Figure: Selecting algorithms and hyperparameters through holdouts

The human as an algorithm.

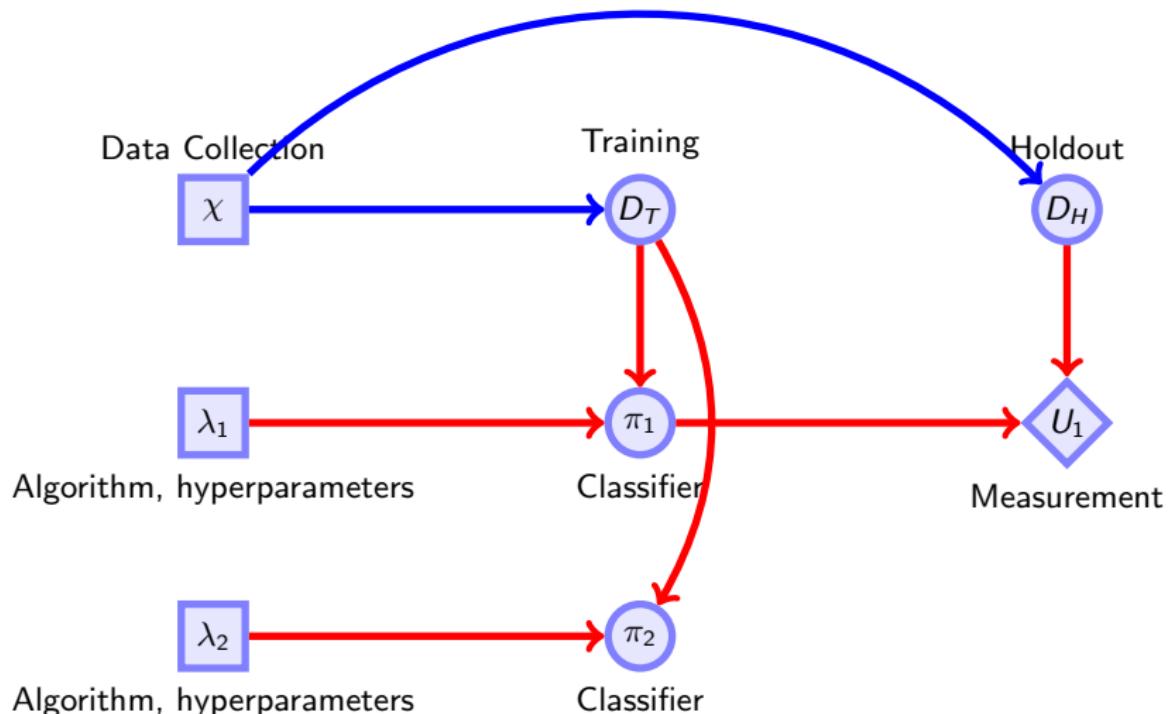


Figure: Selecting algorithms and hyperparameters through holdouts

The human as an algorithm.

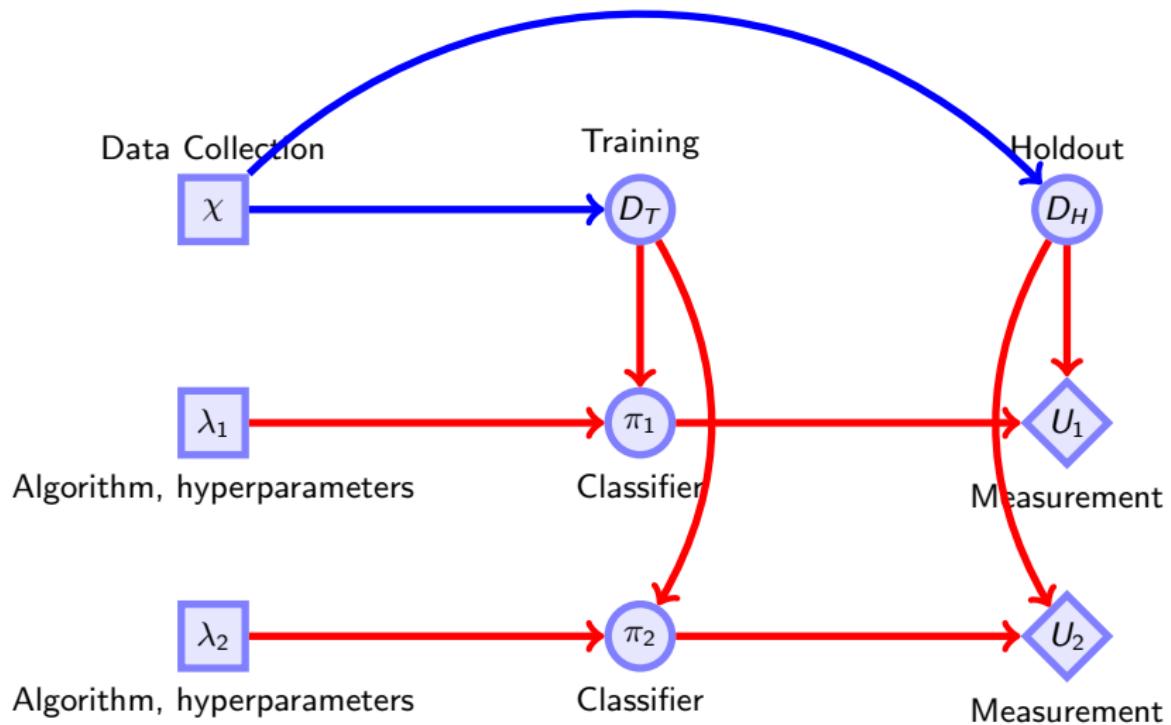


Figure: Selecting algorithms and hyperparameters through holdouts

Holdout sets

- Original data D , e.g. $D = (x_1, \dots, x_T)$.
- Training data $D_T \subset D$, e.g. $D_T = x_1, \dots, x_n$, $n < T$.
- Holdout data $D_H = D \setminus D_T$, used to measure the quality of the result.
- Algorithm λ with hyperparameters ϕ .
- Get algorithm output $\pi = \lambda(D_T, \phi)$.
- Calculate quality of output $U(\pi, D_H)$

Holdout and test sets for unbiased algorithm comparison

Algorithm 2 Unbiased adaptive evaluation through data partitioning

Partition data into D_T, D_H, D^* .

for $\lambda \in \Lambda$ **do**

for $\phi \in \Phi_\lambda$ **do**

$\pi_{\phi, \lambda} = \lambda(D_T, \phi)$.

end for

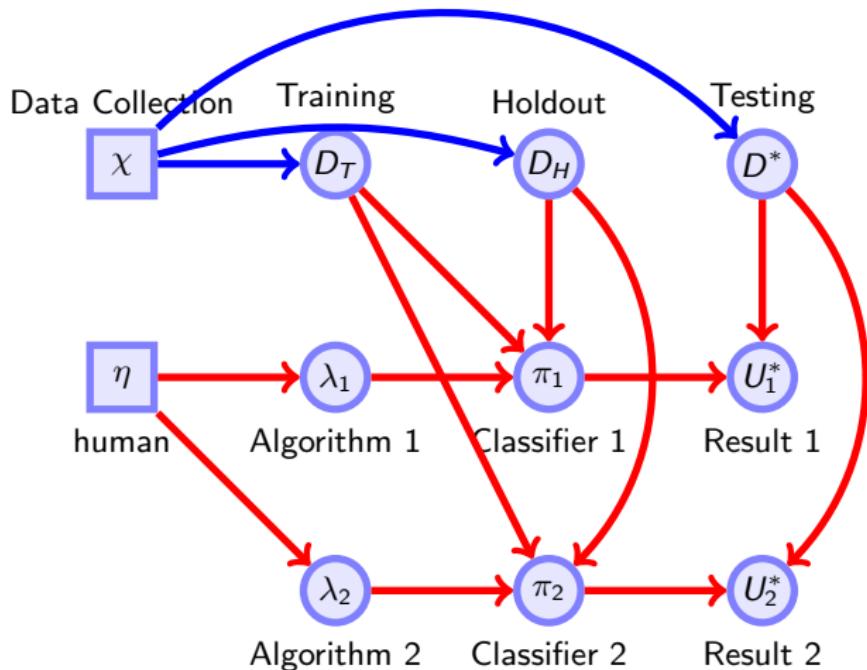
 Get π_λ^* maximising $U(\pi_{\phi, \lambda}, D_H)$.

$u_\lambda = U(\pi_\lambda^*, D^*)$.

end for

$\lambda^* = \arg \max_\lambda u_\lambda$.

Final performance measurement



Independent data sets

Experiment

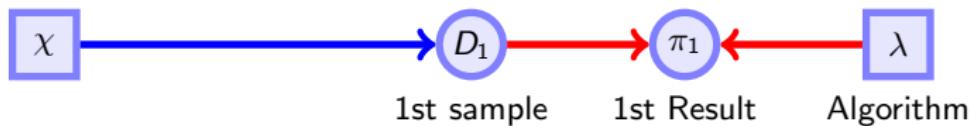


Figure: Multiple samples

Independent data sets

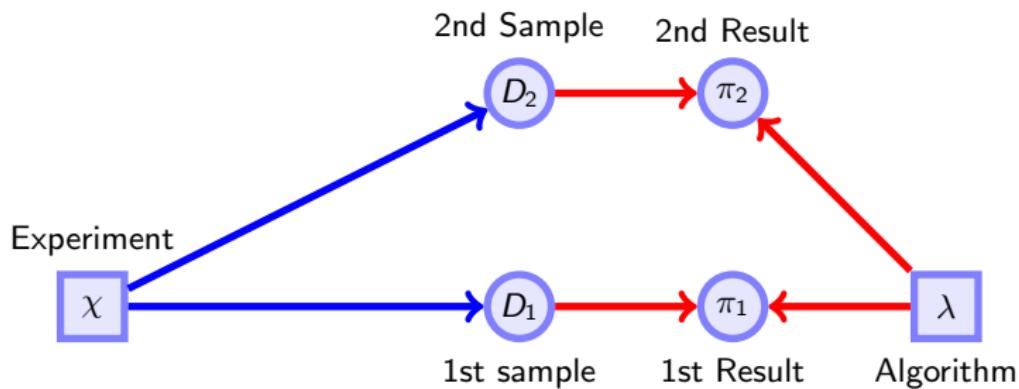


Figure: Multiple samples

Bootstrap samples

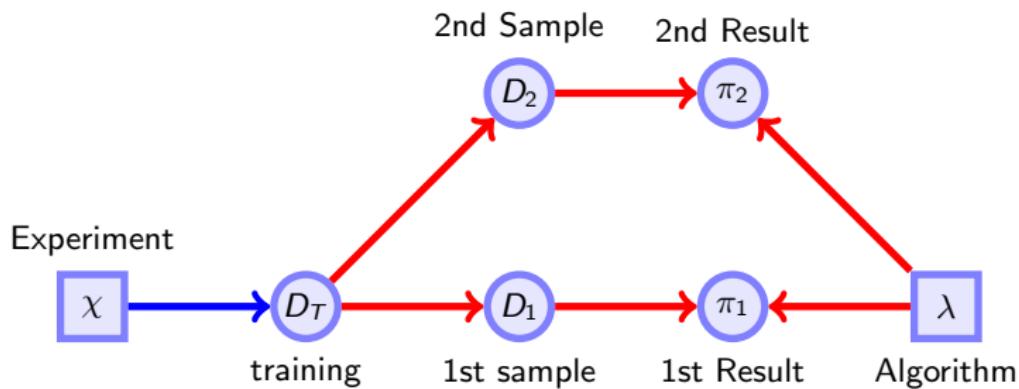


Figure: Bootstrap replicates of a single sample

Bootstrapping

Bootstrapping is a general technique that can be used to:

- Estimate the sensitivity of λ to the data x .
- Obtain a distribution of estimates π from λ and the data x .
- When estimating the performance of an algorithm on a small dataset D^* , use bootstrap samples of D^* .

Bootstrapping

- 1 **Input** Training data D , number of samples k .
- 2 **For** $i = 1, \dots, k$
- 3 $D^{(i)} = \text{Bootstrap}(D)$
- 4 **return** $\{D^{(i)} \mid i = 1, \dots, k\}$.

where $\text{Bootstrap}(D)$ samples with replacement $|D|$ points from D_T .

Cross-validation

k -fold Cross-Validation

- 1 **Input** Training data D_T , number of folds k , algorithm λ , measurement function U
- 2 Create the partition $D^{(1)}, \dots, D^{(k)}$ so that $\bigcup_{i=1}^k D^{(k)} = D$.
- 3 Define $D_T^{(i)} = D \setminus D^{(i)}$
- 4 $\pi_i = \lambda(D_T^{(i)})$
- 5 **For** $i = 1, \dots, k$:
- 6 $\pi_i = \lambda(D^{(i)})$
- 7 $u_i = U(\pi_i)$
- 8 **return** $\{y_1, \dots, y_k\}$.

Independent replication

Replication study

- 1 Reinterpret the original hypothesis and experiment.
- 2 Collect data according to the original protocol, **unless flawed**.
- 3 Run the analysis again, **unless flawed**.
- 4 See if the conclusions are in agreement.

Learning outcomes

Understanding

- What is a hold-out set, cross-validation and bootstrapping.
- The idea of not reusing data input to an algorithm to evaluate it.
- The fact that algorithms can be implemented by both humans and machines.

Skills

- Use git and notebooks to document your work.
- Use hold-out sets or cross-validation to compare parameters/algorithms in Python.
- Use bootstrapping to get estimates of uncertainty in Python.

Reflection

- What is a good use case for cross-validation over hold-out sets?
- When is it a good idea to use bootstrapping?
- How can we use the above techniques to avoid the false discovery problem?
- Can these techniques fully replace independent replication?

- [1] Evelyn Fix and Joseph L Hodges Jr. Discriminatory analysis-nonparametric discrimination: consistency properties. Technical report, California Univ Berkeley, 1951.